

Knowledge Injection in RUL

Let's see the last approach in action

Domain Knowledge as Constraints

Our soft constraint is in the form:

$$f(x_i; \theta) - f(x_j; \theta) = j - i \quad \forall i, j = 1..m \text{ s.t. } c_i = c_j$$

Domain Knowledge as Constraints

Our soft constraint is in the form:

$$f(x_i; \theta) - f(x_j; \theta) = j - i \quad \forall i, j = 1..m \text{ s.t. } c_i = c_j$$

From which we can derive the penalizer:

$$\lambda \sum_{\substack{i, j=1..m \\ c_i=c_j}} \left(f(x_i; \theta) - f(x_j; \theta) - (j - i) \right)^2$$

- For sake of simplicity, we are using the same λ for all pairs
- ...Meaning the same standard deviation using a Normal semantic

Domain Knowledge as Constraints

Our soft constraint is in the form:

$$f(x_i; \theta) - f(x_j; \theta) = j - i \quad \forall i, j = 1..m \text{ s.t. } c_i = c_j$$

From which we can derive the penalizer:

$$\lambda \sum_{\substack{i, j=1..m \\ c_i=c_j}} \left(f(x_i; \theta) - f(x_j; \theta) - (j - i) \right)^2$$

- For sake of simplicity, we are using the same λ for all pairs
- ...Meaning the same standard deviation using a Normal semantic

There's large number of constraints, but many are redundant

Our Regularizer

For example, we can focus on subsequent pairs

$$L(\hat{y}) + \lambda \sum_{\substack{i < j \\ c_i = c_j}} \left(f(x_i; \theta) - f(x_j; \theta) - (j - i) \right)^2$$

- Where $i < j$ iff j is the next sample for after i for a given machine
- This approach requires a linear (rather than quadratic) number of constraints

This method can work with mini-batches

- In this case, $<$ will refer to contiguous samples in the same batch
- ...And of course for the same component

We will now see how to implement this approach

Generating Batches from the Same Machine

Our regularizer requires to have **sorted** samples **from the same machine**

The easiest way to ensure we have enough is using a custom `DataGenerator`

```
class SMBatchGenerator(tf.keras.utils.Sequence):  
    def __init__(self, data, in_cols, batch_size, seed=42): ...  
    def __len__(self): ...  
    def __getitem__(self, index): ...  
    def on_epoch_end(self): ...  
    def __build_batches(self): ...
```

- `__len__` is called to know how many batches are left
- `__getitem__` should return one batch
- `on_epoch_end` should take care (e.g.) of shuffling

Generating Batches from the Same Machine

The `__init__` method takes care of the initial setup

```
def __init__(self, data, in_cols, batch_size, seed=42):  
    super(SMBatchGenerator).__init__()  
    self.data = data  
    self.in_cols = in_cols  
    self.dpm = split_by_field(data, 'machine')  
    self.rng = np.random.default_rng(seed)  
    self.batch_size = batch_size  
    # Build the first sequence of batches  
    self.__build_batches()
```

- We store some fields
- We split the data by machine
- We build a dedicated RNG
- ...And finally we call the custom-made `__build_batches` method

Generating Batches from the Same Machine

The `__build_batches` method prepares the batches for one full epoch

```
def __build_batches(self):
    self.batches, self.machines = [], []
    mcns = list(self.dpm.keys())
    self.rng.shuffle(mcns) # sort the machines at random
    for mcn in mcns: # Loop over all machines
        index = self.dpm[mcn].index # sample indexes for this machine
        ...
        self.rng.shuffle(idx) # shuffle sample indexes for this machine
        bt = idx.reshape(-1, self.batch_size) # split into batches
        bt = np.sort(bt, axis=1) # sort every batch individually
        self.batches.append(bt) # store the batch
        self.machines.append(np.repeat([mcn], len(bt))) # add machine information
    self.batches = np.vstack(self.batches) # concatenate
    self.machines = np.hstack(self.machines)
```


Generating Batches from the Same Machine

We rebuild batches after each epoch

```
def on_epoch_end(self):  
    self.__build_batches()
```

Most of the remaining work is done in the `__getitem__` method:

```
def __getitem__(self, index):  
    idx = self.batches[index]  
    x = self.data[self.in_cols].loc[idx].values  
    y = self.data['rul'].loc[idx].values  
    flags = (y != -1)  
    info = np.vstack((y, flags, idx)).T  
    return x, info
```

- The RUL value is -1 for the unsupervised data: we flag the meaningful RULs
- ...We pack indexes, RUL values, and flags into a single **info** tensor

Custom Training Step

We then enforce the constraints by means of a **custom training step**

```
class CstRULRegressor(keras.Model):  
    def __init__(self, rul_pred, alpha, beta, maxrul): ...  
  
    def train_step(self, data): ...  
  
    def call(self, data): return self.rul_pred(data)  
  
...
```

- We use a custom `keras.Model` subclass
- ...And accept an externally built RUL prediction model (`rul_pred`)
- The custom training step is implemented in `train_step`
- The `call` method relies on the external model for RUL prediction

Custom Training Step

In the `__init__` function:

```
def __init__(self, rul_pred, alpha, beta, maxrul):  
    super(CstRULRegressor, self).__init__(input_shape, hidden)  
    # Store the base RUL prediction model  
    self.rul_pred = rul_pred  
    # Weights  
    self.alpha = alpha  
    self.beta = beta  
    self.maxrul = maxrul  
    ...
```

- **beta** is the regularizer weight, **alpha** is a weight for the loss function itself
- We also store the maximum RUL

Custom Training Step (CHANGE THIS)

In the custom training step:

```
def train_step(self, data):
    x, info = data
    y_true, flags, idx = info[:, 0:1], info[:, 1:2], info[:, 2:3]
    with tf.GradientTape() as tape:
        y_pred = self(x, training=True) # predictions
        mse = k.mean(flags * k.square(y_pred - y_true)) # MSE loss
        delta_pred = y_pred[1:] - y_pred[:-1] # pred. difference
        delta_rul = -(idx[1:] - idx[:-1]) / self.maxrul # index difference
        deltadiff = delta_pred - delta_rul # difference of differences
        cst = k.mean(k.square(deltadiff)) # regularization term
        loss = self.alpha * mse + self.beta * cst # loss
    ...
```

- We unpack the `info` tensor
- Inside a `GradientTape`, we construct our regularized loss

Custom Training Step

In the custom training step:

```
def train_step(self, data):  
    ...  
    tr_vars = self.trainable_variables  
    grads = tape.gradient(loss, tr_vars) # gradient computation  
  
    self.optimizer.apply_gradients(zip(grads, tr_vars)) # weight update  
  
    ...
```

- We then apply the (Stochastic) Gradient Descent step
- Then we update and return the loss trackers

Training the Lagrangian Approach

We can now test our approach


```
In [3]: nn_aux = util.build_ml_model(input_size=len(dt_in), output_size=1, hidden=[32, 32])
nn2 = util.CstRULRegressor(rul_pred=nn_aux, alpha=1, beta=5, maxrul=maxrul)
batch_gen = util.CstBatchGenerator(tr_s2, dt_in, batch_size=32)
history = util.train_ml_model(nn2, X=batch_gen, y=None, validation_split=0., epochs=25, verl
```

```
Epoch 1/25
1071/1071 _____ 2s 1ms/step - cst: 0.0048 - loss: 0.0408 - mse: 0.0167
Epoch 2/25
1071/1071 _____ 2s 1ms/step - cst: 0.0020 - loss: 0.0123 - mse: 0.0023
Epoch 3/25
1071/1071 _____ 2s 1ms/step - cst: 0.0021 - loss: 0.0129 - mse: 0.0026
Epoch 4/25
1071/1071 _____ 1s 1ms/step - cst: 0.0019 - loss: 0.0118 - mse: 0.0023
Epoch 5/25
1071/1071 _____ 1s 1ms/step - cst: 0.0019 - loss: 0.0123 - mse: 0.0025
Epoch 6/25
1071/1071 _____ 2s 1ms/step - cst: 0.0019 - loss: 0.0123 - mse: 0.0027
Epoch 7/25
1071/1071 _____ 2s 1ms/step - cst: 0.0019 - loss: 0.0125 - mse: 0.0028
Epoch 8/25
1071/1071 _____ 2s 1ms/step - cst: 0.0019 - loss: 0.0122 - mse: 0.0027
Epoch 9/25
1071/1071 _____ 2s 1ms/step - cst: 0.0018 - loss: 0.0117 - mse: 0.0025
Epoch 10/25
1071/1071 _____ 1s 1ms/step - cst: 0.0018 - loss: 0.0116 - mse: 0.0026
Epoch 11/25
1071/1071 _____ 1s 1ms/step - cst: 0.0018 - loss: 0.0116 - mse: 0.0026
```

Training the Lagrangian Approach

...And we can check the training curve

```
In [15]: util.plot_training_history(history, figsize=figsize)
```




Model loss: 0.0110 (training)

Inspecting the Predictions

Let's have a look at the predictions on the **test** data

```
In [16]: ts_pred2 = nn2.predict(ts_s[dt_in], verbose=0).ravel() * maxrul  
util.plot_rul(ts_pred2[:stop], ts["rul"].iloc[:stop], same_scale=False, figsize=figsize)
```




- The signal is **much more stable**
- The scale is still off, but we can fix that with a well chosen threshold

Threshold Optimization and Cost Evaluation

We can now optimize the threshold optimization (on the supervised data)

```
In [19]: cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=safe_interval)
th_range2 = np.linspace(70, 150, 200)
trs_thr2 = util.optimize_threshold(trs_s['machine'].values, trs_pred2, th_range2, cmodel, p
print(f'Optimal threshold for the training set: {trs_thr2:.2f}')
```

Optimal threshold for the training set: 98.14



Threshold Optimization and Cost Evaluation

Finally, we can evaluate the SBR approach in terms of cost

```
In [21]: trs_c2, trs_f2, trs_sl2 = cmodel.cost(trs_s['machine'].values, trs_pred2, trs_thr2, return_  
tru_c2, tru_f2, tru_sl2 = cmodel.cost(tru_s['machine'].values, tru_pred2, trs_thr2, return_  
ts_c2, ts_f2, ts_sl2 = cmodel.cost(ts['machine'].values, ts_pred2, trs_thr2, return_margin=  
print(f'Cost: {trs_c2/len(trs_mcn):.2f} (supervised), {tru_c2/len(tru_mcn):.2f} (unsupervised)
```

Cost: -36.86 (supervised), -72.11 (unsupervised), -79.34 (test)

```
In [22]: print(f'Avg. fails: {trs_f2/len(trs_mcn):.2f} (supervised), {tru_f2/len(tru_mcn):.2f} (unsup  
print(f'Avg. slack: {trs_sl2/len(trs_mcn):.2f} (supervised), {tru_sl2/len(tru_mcn):.2f} (uns
```

Avg. fails: 0.00 (supervised), 0.02 (unsupervised), 0.00 (test)

Avg. slack: 32.14 (supervised), 37.84 (unsupervised), 35.28 (test)

- The number of fails has decreased very significantly
- The slack is still contained

And we did this with **just a handful** of run-to-failure experiments