

# Constrained ML via Inference-Time Projection

---

It's never too late to quit violating constraints

# From Data Projection

**Let's consider again our data projection approach**

First we project the training data:

$$z^* = \operatorname{argmin}_z \{ L(z, y) \mid g(z) \leq 0 \}$$

...Then we train a supervised learning model as usual:

$$\theta^* = \operatorname{argmin}_{\theta} \{ L(\hat{y}, z^*) \mid \hat{y} = f(x; \theta) \}$$

**By doing this, we push the model towards learning the constraints**

...But we do it by decomposing the original problem

- This leads to some significant advantages (mainly flexibility)
- ...And a few disadvantages (e.g. less precise approximation)

## ...To Inference-Time Projection

**What if we swapped the two steps?**

We could start by training a model as usual:

$$\theta^* = \operatorname{argmin}_{\theta} \{ L(\hat{y}, y) \mid \hat{y} = f(x; \theta) \}$$

...And then projecting the predictions:

$$\hat{z} = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

## ...To Inference-Time Projection

### What if we swapped the two steps?

We could start by training a model as usual:

$$\theta^* = \operatorname{argmin}_{\theta} \{ L(\hat{y}, y) \mid \hat{y} = f(x; \theta) \}$$

...And then projecting the predictions:

$$\hat{z} = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

- We still perform projection in the output space
- ...But in this case we do it at inference-time

**For this reason, we'll call this approach inference-time projection**

# Probabilistic Semantic

As for data projection, when  $L$  represents a likelihood

...We can see the projection process as MAP computation:

$$\hat{z} = \underset{z}{\operatorname{argmin}} \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

- We seek a feasible vector  $\hat{z}$
- ...That has the **largest estimated probability** w.r.t. the predictions

**This trick is extensively used in ML!**

- Rounding in binary classification is a form of projection
- ...And the same goes for the **argmax** used in multiclass classification

**Let's start to study in detail this approach**

## Two is Complicated, Three is a Crisis

To begin with, our original does not really work

$$\hat{z} = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

This projection problem is well defined for the training data

- However, since we are projecting at inference time
- ...We'll need to do it for unseen examples as well

# Two is Complicated, Three is a Crisis

To begin with, our original does not really work

$$\hat{z} = \underset{z}{\operatorname{argmin}} \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

This projection problem is well defined for the training data

- However, since we are projecting at inference time
- ...We'll need to do it for unseen examples as well

## Unseen examples might not come in batches

...And even when they do, they may not be representative samples

- So, it's hard to deal with relational constraints and distribution constraints
- In most cases, we'll be projecting one example at a time

...Though there are exceptions to this rule



## Where did the Input Go?

**All our constraint formulations so far make no mention of the input**

For example, in our projection problem

$$\hat{z} = \underset{z}{\operatorname{argmin}} \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

...The constraint function  $g(x)$  is defined on the output space alone

## Where did the Input Go?

**All our constraint formulations so far make no mention of the input**

For example, in our projection problem

$$\hat{z} = \underset{z}{\operatorname{argmin}} \{ L(z, f(x; \theta^*)) \mid g(z) \leq 0 \}$$

...The constraint function  $g(x)$  is defined on the output space alone

**When working with the training data, this is totally fine**

- It's quite frequent for the input to have an impact on the constraint
- ...But the input is **fixed** at training time
- ...And therefore we can include its effect in the definition of  $g$

# Where did the Input Go?

**On unseen examples, this is no longer the case**

...And we need to account for the input explicitly:

$$\hat{z}(x) = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(x, z) \leq 0 \}$$

Formally:

- The projection  $\hat{z}$  will be always input-dependent
- The constraint function  $g(x, y)$  will be input-depended in general

# Where did the Input Go?

**On unseen examples, this is no longer the case**

...And we need to account for the input explicitly:

$$\hat{z}(x) = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(x, z) \leq 0 \}$$

Formally:

- The projection  $\hat{z}$  will be always input-dependent
- The constraint function  $g(x, y)$  will be input-depended in general

**In practice, this isn't a big change**

- We'll need to build our constraints based on the value of  $x$
- ...And then we can solve the projection problem as usual

# Computational Effort

The computational effort for our projection operator

$$\hat{z}(x) = \operatorname{argmin}_z \{ L(z, f(x; \theta^*)) \mid g(x, z) \leq 0 \}$$

...Is similar to the one we had for data-projection

- The fact that we typically project individual examples is good news
- ...Since it makes our problems **smaller** and **simpler**

# Computational Effort

The computational effort for our projection operator

$$\hat{z}(x) = \underset{z}{\operatorname{argmin}} \{ L(z, f(x; \theta^*)) \mid g(x, z) \leq 0 \}$$

...Is similar to the one we had for data-projection

- The fact that we typically project individual examples is good news
- ...Since it makes our problems **smaller** and **simpler**

**However, there are cases when even this can be an issue**

Some applications (e.g. control systems) require very low latency

- In this case, we might not have time enough to project

Some hardware has trouble with projection

- E.g. no optimization packages or not enough power on cyber-physical systems

# The Bad, the Ugly, and the Good

**So far, we've focused on the issues of this approach**

- It's hard to deal with many examples, which is bad
- There's a non-trivial computation step at inference time, which is ugly

So, why should we use inference-time projection at all?

# The Bad, the Ugly, and the Good

**So far, we've focused on the issues of this approach**

- It's hard to deal with many examples, which is bad
- There's a non-trivial computation step at inference time, which is ugly

So, why should we use inference-time projection at all?

Three words.

**Guaranteed Constraint Satisfaction**

...And it holds also out-of-distribution!

- In some cases this can be a massive advantage
- Inference time projection is the simplest method to achieve it



**Let's consider a simple case study for inference-time  
projection**

# Mutually Exclusive Classes

Let's consider a **multilabel classification** problem

- Given an input sample
- ...We need to assign to it zero-to- $k$  classes

**The problem arises in a number of settings**


- Assigning tags to posts, or items in a marketplace
- Estimating properties of materials
- Identifying the favorite music/movie genres for a user
- ...

**We can build individual classifiers for each class**

...But using a single model for them all allows to exploit **correlations**

# Yeast Traits

We'll focus on estimating genetic traits of yeast cells



- Given numeric data from a micro-array analysis
- ...We want to estimate the characteristic of a given yeast

# Loading the Data

## Let's load our dataset

...Which comes from the [Multi-label Classification Dataset Repository](#)

```
In [85]: fname = os.path.join '..', 'data', 'Yeast.arff')
X_tr, X_ts, y_tr, y_ts = util.load_multilabel_dataset(fname, label_num=14, standardize=True)
display(y_tr.head())
```


	Class1	Class2	Class3	Class4	Class5	Class6	Class7	Class8	Class9	Class10	Class11	Class12	Class13	Class14
<b>1681</b>	0	1	1	1	0	0	0	0	0	1	1	1	1	0
<b>1565</b>	1	1	0	0	0	1	1	0	0	0	0	1	1	0
<b>544</b>	0	1	1	1	1	1	0	0	0	0	0	0	0	0
<b>1517</b>	0	0	0	0	0	0	1	1	0	0	0	1	1	0
<b>821</b>	1	1	0	0	0	0	0	0	0	0	0	1	1	0

- All inputs are numeric
- ...And we have 14 possible classes

# Unconstrained Model

We'll start by training an unconstrained model

```
In [86]: nn0 = util.build_ml_model(input_size=X_tr.shape[1], output_size=y_tr.shape[1], hidden=[32],  
history = util.train_ml_model(nn0, X_tr, y_tr, epochs=200, validation_split=0.1, loss='binary_crossentropy',  
util.plot_training_history(history, figsize=figsize))
```




Model loss: 0.3658 (training) 0.4431 (validation)

# Accuracy Evaluation

We can evaluate the model in terms of per-class accuracy

```
In [87]: p_ts = nn0.predict(X_ts, verbose=0)
c_ts = np.round(p_ts)
ts_acc = util.get_multi_label_accuracy(y_ts, c_ts, series_name='original')
ts_acc.plot.bar(figsize=figsize, xlabel='class', ylabel='accuracy');
print(f'Average class accuracy: {ts_acc.mean():.3f} (test)')
```

Average class accuracy: 0.802 (test)



# Mutually Exclusive Classes

**We will assume that certain traits cannot occur together**

...Meaning that certain classes are mutually exclusive

- In this case, the constraint is just a purely synthetic example
- In the real world, constraints may come from regulation, domain knowledge, etc.

**We will treat as mutually exclusive all pairs of classes**

...That have a low-enough rate of co-occurrence **in the entire dataset**


```
In [88]: all_y = pd.concat((y_tr, y_ts))  
         forbidden = util.get_forbidden_pairs(all_y, cutoff_quantile=0.5)  
         print(forbidden)
```

```
{(4, 9), (3, 7), (4, 6), (12, 13), (3, 10), (5, 13), (3, 13), (5, 10), (8, 9), (8, 12),  
(1, 6), (0, 8), (2, 5), (2, 8), (7, 10), (6, 8), (7, 13), (3, 9), (4, 8), (3, 6), (5, 9),  
(8, 11), (0, 7), (0, 4), (0, 10), (2, 7), (9, 13), (0, 13), (2, 10), (1, 8), (11, 13), (2,  
13), (7, 9), (6, 13), (6, 10), (3, 5), (4, 10), (3, 8), (4, 13), (5, 8), (0, 3), (0, 9),  
(8, 10), (10, 13), (0, 6), (8, 13), (1, 13), (2, 6), (6, 9)}
```

# Checking Violations

We can check how often our constraints are violated on the test data

```
In [100]: util.forbidden_pair_violation(c_ts, forbidden).plot.bar(figsize=figsize, ylabel='violation rate')
```



There some violation, though not many



# Projection Model

We now need to define our projection problem:

$$\hat{z}(x) = \underset{z}{\operatorname{argmin}} \{ L(z, \hat{y}) \mid g(x, z) \leq 0 \} \quad \text{with: } \hat{y} = f(x; \theta^*)$$

For pre-computed  $\hat{y}$ , we'll need to obtain a declarative model of:

- The decision variables  $z$
- The loss function  $L(z, \hat{y})$
- The constraints  $g(x, z) \leq 0$

# Projection Model

We now need to define our projection problem:

$$\hat{z}(x) = \underset{z}{\operatorname{argmin}} \{ L(z, \hat{y}) \mid g(x, z) \leq 0 \} \quad \text{with: } \hat{y} = f(x; \theta^*)$$

For **pre-computed**  $\hat{y}$ , we'll need to obtain a declarative model of:

- The decision variables  $z$
- The loss function  $L(z, \hat{y})$
- The constraints  $g(x, z) \leq 0$

The decision variables represent a **vector of classes**

$$z_i \in \{0, 1\} \quad \text{with } z_i = 1 \text{ iff we predict class } i$$

# Projection Model

**We'll use the binary cross entropy as our loss function**

- The decision variables represent **actual classes**
- ...While the  $\hat{y}$  contains estimated probabilities

**Hence, we can model the loss function as:**

$$L(z, \hat{y}) = - \sum_{i=1}^k z_i \log(\hat{y}_i + \epsilon) + (1 - z_i) \log(1 - \hat{y}_i + \epsilon)$$

- Since the cross-entropy is not well defined for  $\hat{y}$  equal to 0 or 1
- ...We add a small constant to the logarithm

Overall, we have a **linear expression** defined on  $z$

# Projection Model

**The constraints state that certain pairs of classes cannot occur together**

...Which can be modeled as:

$$z_i + z_j \leq 1 \quad \forall (i, j) \in F$$

- We consider all pairs of mutually exclusive classes
- ...And we allow at most one of the associated  $z$  variables to be 1

**Overall, we have the following Mixed Integer Linear Program**

$$\underset{z}{\operatorname{argmin}} z^T \log(\hat{y} + \varepsilon) + (1 - z)^T \log(1 - \hat{y} + \varepsilon)$$

$$\text{subject to: } z_i + z_j \leq 1$$

$$z_i \in \{0, 1\}$$

$$\forall (j, j) \in F$$


$$\forall i = 1..n_c$$

# Maximum Weight Independent Set

What we are looking for is a maximum weight independent set

...On the graph defined by the mutually exclusive classes

```
In [101]: util.plot_forbidden_pairs(nclasses=y_tr.shape[1], forbidden=forbidden, figsize=figsize)
```



- This is an NP-hard problem, but **not a very difficult one!**

# Implementation

An implementation of the model is available in the `util` module

```
def project_multilabel_classification(p, forbidden, solver='CBC', tlim=10):
    ...
    z = [slv.IntVar(0, 1, 'z_{c}') for c in range(label_num)]
    for c0, c1 in forbidden: # constraints
        slv.Add(z[c0] + z[c1] <= 1)
    for ii in range(p.shape[0]):
        # Define the objective
        obj = - sum([z[c] * np.log(p[ii, c] + 1e-4) +
                     (1 - z[c]) * np.log(1 - p[ii, c] + 1e-4)
                     for c in range(label_num)])
        slv.Minimize(obj)
        status = slv.Solve()
    ...
```

- For improved efficiency, we build a single model to project multiple examples
- On each example, we just change the objective function

# Computing the Projections

We can now compute the inference-time projections for the test set

```
In [104]: %time pc_ts = util.project_multilabel_classification(p_ts, forbidden=forbidden, tlim=10, so
print(f'For {len(p_ts)} test examples')
```

```
CPU times: user 194 ms, sys: 31.6 ms, total: 225 ms
Wall time: 225 ms
For 484 test examples
```

- Even if the problem is technically intractable (NP-hard)
- ...The run time is **very low**!

For roughly 10% of the example the projection results in a change:


```
In [105]: nc_ts = np.any(np.abs(pc_ts - c_ts) > 0, axis=1).sum()
print(f'Non-trivial projections: {nc_ts} (test)')
```

```
Non-trivial projections: 44 (test)
```

# Accuracy Evaluation

We can compare the accuracy of the constrained model with the original one

```
In [109]: tmp = pd.concat([ts_acc_p, ts_acc], axis=1)
tmp.plot.bar(figsize=(figsize[0], figsize[1]*0.8), xlabel='class', ylabel='accuracy');
```






# Accuracy Evaluation

We can compare the accuracy of the constrained model with the original one

```
In [109]: tmp = pd.concat([ts_acc_p, ts_acc], axis=1)
tmp.plot.bar(figsize=(figsize[0], figsize[1]*0.8), xlabel='class', ylabel='accuracy');
```




- For some classes, the constrained model works even better!
- This happens since the constraint encode some information about unseen examples

# Accuracy Evaluation

Projection methods enable also a comparison **against an oracle**

Which can be obtained by **projecting the ground truth labels**

```
In [112]: ts_acc_o = util.get_multi_label_accuracy(y_ts, oc_ts, series_name='oracle')
tmp = pd.concat([ts_acc_p, ts_acc, ts_acc_o], axis=1)
tmp.plot.bar(figsize=(figsize[0], figsize[1]*0.7), xlabel='class', ylabel='accuracy');
```



Doing this tells us **the best we can do**, accounting for the constraints