

# Ordinary Differential Equations

---



# Ordinary Differential Equations

An **Ordinary Differential Equation** is any equation in the form:

$$\dot{y} = f(y, t)$$

- Where  $y$  is the **state variable**
- ...And  $f$  is a function, providing the gradient of the state variable

## The peculiarities:

- $y$  is actually a **function** of the  $t$  variable
- The  $t$  variable typically (but not always) represents **time**
- ...Hence  $y(t)$  is the state at time  $t$
- The gradient  $f$  depends on both the current state and current time

**Ordinary** = does not feature partial derivatives



# Initial Value Problem

An **Initial Value Problem** consists of an ODE and a initial condition

$$\begin{aligned}\dot{y} &= f(y, t) \\ y(0) &= y_0\end{aligned}$$

- This can be interpreted as **running a simulation**
- Given that the initial state  $y(0)$  is  $y_0$ , how will the state unfold?

**Initial values problem can be solved (a.k.a. integrated):**

Exactly, using **symbolic approaches**, e.g.

$$\dot{y} = a, y(0) = b \quad \Rightarrow \quad y(t) = ay + b$$


- This is the method considered in typical calculus courses

...Or approximately, via **numerical approach**



## An Example

As an example, let's consider a simple RC circuit



It's dynamic behavior is described by the ODE:

$$\dot{V} = \frac{1}{\tau}(V_s - V)$$

- Where  $\tau = RC$



# Euler Method

The simplest numerical approach for ODEs is called **Euler Method**

This is obtained by:

- Considering a fixed sequence of evaluation points  $\{t_k\}_{k=0}^n$
- Using a **linear approximation** for  $y(t)$  within each interval  $[t_k, t_{k+1}]$
- Approximating the slope with the gradient at time  $t_k$

The pseudo code of the method consists of **a single loop**

- for  $k = 1..n$ :
  - $y_k = y_{k-1} + (t_k - t_{k-1})f(y_{k-1}, t_{k-1})$

The output is a sequence  $\{y_k\}_{i=0}^n$

- $y_k$  is the state at time  $t_k$
- $y_0$  is also an input for the algorithm



# Euler Method for the RC Circuit

A typical Initial Value Problem solver API requires to define

The function characterizing the equation, i.e.  $f(y, t)$ :

```
In [2]: tau, Vs = 8, 12  
f = lambda y, t: 1./tau * (Vs - y)
```

The initial state  $y_0$  and the evaluation points  $\{t_i\}_{i=0}^n$

```
In [3]: y0 = (0,) # We start from an empty capacitor  
t = np.linspace(0, 40, 12)
```

Then we can call the solver itself (the code is in the `util` module)


```
In [4]: y, dy = util.euler_method(f, y0, t, return_gradients=True)
```



# Euler Method for the RC Circuit

Visually, the method works as follows:

```
In [5]: util.plot_euler_method(y, t, dy=dy, xlabel='time', ylabel='V', figsize=figsize)
```



- The dots represent evaluated states



- The slope of the lines corresponds to the gradient at each step

# ODE Integration Methods

## Euler method is the simplest ODE integration approach

...But also one of the worst in terms of accuracy

- This is due to errors in the local approximation
- ...And forces to use very small steps to obtain high-quality results

## There are many alternative integration methods

Some examples include:

- Backward Euler method
  - Like Euler method, but we use the gradient at the **next** state
  - In practice it requires to solve a (typically non-linear) equation
- Runge-Kutta methods
  - It's a family of method (Euler method is the simplest version)
  - They combine multiple gradients to obtain a local slope





# "Learning" ODEs

---



# "Learning" ODEs

The parameters of an ODE can be **estimated** from data

Formally, training amounts to solving:

$$\operatorname{argmin}_{\omega} \left\{ L(\hat{y}(t), y) \mid \dot{\hat{y}} = f(\hat{y}, t; \theta), \hat{y}(0) = y_0 \right\}$$

Where:

- $\{t_k\}_{k=0}^n$  is a sequence of points for which measurements are available
- $\{y_k\}_{k=0}^n$  are the corresponding state measurements
- $f$  is a **parameterized** gradient function
- $L$  is a loss function (e.g. the classical MSE)

Intuitively, we require the integrated ODE to be close to the real one

- The goal is to choose the parameters (e.g.  $\tau, V_s$ ) so as to achieve this



# "Learning" ODEs

**A viable approach is to "discretize, then optimize"**

...Which can be done by relying on an automatic differentiation engine

- First we solve the initial value problem using a numerical method
  - ...Making sure we evaluate every point in  $\{t_k\}$
- Then, we compute the loss  $L$
- ...And view the whole process as a single compute graph

Then optimization over  $\theta$  can be performed (e.g.) via gradient descent

**This is possible since every integration step is differentiable**

This is true for Euler method, but also for other (better) integration methods

- In particular, it's doable for the whole Runge-Kutta family
- ...But a bit more complicated in implicit methods (e.g. backward Euler)



# Building Our Ground Truth

## We'll see an example using our simple RC circuit

Let's start by building a high-quality ground truth sequence

- We will use the `odeint` solver from scikit learn for this
- The code can be found in the `simulate_RC` function

```
In [6]: V0, tau, Vs, tmax = 0, tau, Vs, 60
data = util.simulate_RC(V0, tau, Vs, tmax, steps_per_unit=1)
data.head()
```

Out[6]:


V	
time	
0.0	0.000000
1.0	1.410037
2.0	2.654391
3.0	3.752529
4.0	4.721632

- `steps_per_unit` defines how many evaluations to perform per unit of time

# Building Our Ground Truth

Let' check (visually) that the result is smooth enough

```
In [7]: util.plot_df_cols(data, figsize=figsize)
```



- We need this step since we treat this as the ground truth

✎ ... Even if it is a numerical approximation

# Outline of the Approach

**We will go for a simple, but relatively general approach**

- We will view the (parameterized) gradient function  $f(y, t; \theta)$  as a **layer type**
- ...And we will use a `keras.Model` to encode Euler method, i.e.

$$\hat{y}(t_k) = \hat{y}(t_{k-1}) + (t_k - t_{k-1})f(\hat{y}(t_{k-1}), t_{k-1}; \theta)$$

- Each step of the method can be viewed as **layer instance**
- ...And all instances share the same weights


**In terms of input/output:**

- The initial state corresponds to the **input**
- ...And a **secondary input** is given by the sequence  $\{t_k\}_{k=0}^n$
- The **output** is the state for each evaluation step



# Outline of the Approach

Overall, our "architecture" looks like this:



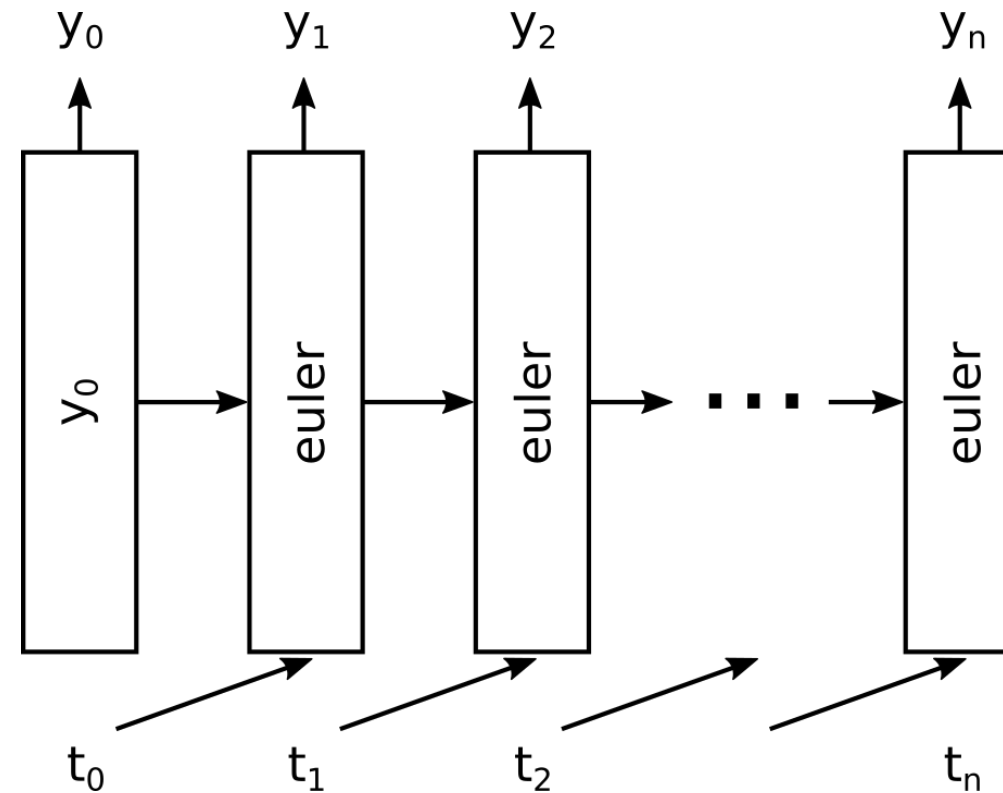
- The input includes the initial state  $y_0$  and the evaluation points  $\{t_k\}_{k=0}^n$
- The output consists of the sequence of state evaluations  $\{\hat{y}_k\}_{k=0}^n$

Overall, the signature is analogous to that of an ODE solver



# Outline of the Approach

Overall, our "architecture" looks like this:



- Each "example" corresponds to a distinct integration of the same system
- ...And the architecture is very similar to a recurrent NN
- In particular, the "depth" grows with the number of evaluation points





## Details Matter

In our RC circuit case, we have:

$$\begin{aligned} & \operatorname{argmin}_{\omega} L(\hat{y}(t), y) \\ & \text{subject to } \dot{\hat{y}} = \frac{1}{\tau}(V_s - \hat{y}) \\ & \hat{y}(0) = y_0 \end{aligned}$$

- Where the parameters to be learned are  $\tau$  and  $V_s$

**There are a few details we need to account for**

- For both parameters, **negative values** make no sense
- Moreover, since we plan to use gradient descent for training
- ...We need to make sure that our initial guesses are **reasonable**



## Details Matter

We can meet both conditions by adopting the reformulation:

$$\tau = \sigma_{\tau} e^{\theta_{\tau}}$$
$$V_s = \sigma_{V_s} e^{\theta_{V_s}}$$

Where the parameters to be learned are now  $\theta_{\tau}$  and  $\theta_{V_s}$

- Using an exponential ensures we get non-negative values
- The scaling factors  $\sigma_{\tau}$  and  $\sigma_{V_s}$  are user-provided
  - They lead to reasonable guesses for typical NN weight initializers

There are just a few mild downsides:

- The exponential may lead to numerical issues in edge cases
- We need to have a rough idea of the scale of  $\tau$  and  $V_s$



# RC Circuit Layer

The layer for the RC circuit gradient is in the `RNablaLayer` class

```
class RNablaLayer(keras.layers.Layer):
    def build(self, input_shape):
        ...
        self.logtau = self.add_weight(
            shape=(1, ),
            initializer = 'random_normal', trainable=True,
        )
        ...
        self.logvs = self.add_weight(
            shape=(1, ),
            initializer = 'random_normal', trainable=True,
        )
        ...
```

- In the `__init__` method we take care of weight initialization



# RC Circuit Layer

The layer for the RC circuit gradient is in the `RNablaLayer` class

```
class RNablaLayer(keras.layers.Layer):  
    ...  
  
    def get_tau(self):  
        return tf.math.exp(self.logtau) * self.tau_ref  
  
    def get_vs(self):  
        return tf.math.exp(self.logvs) * self.vs_ref  
  
    def call(self, inputs):  
        y, t = inputs # unpack the inputs  
        return 1. / self.get_tau() * (self.get_vs() - y)
```

- We use dedicated method to obtain  $\tau$  and  $V_s$
- In the `call` method we compute the (ODE) gradient



# Euler Method Model

The model for the Euler method is in the `ODEEulerModel` class

```
class ODEEulerModel(keras.Model):
    def __init__(self, f, **params): ...

    def call(self, inputs, training=False):
        y, T = inputs # unpack
        res = [y] # initial state
        for i in range(T.shape[1]-1):
            t, nt = T[:, i:i+1], T[:, i+1:i+2] # t_k and t_{k+1}
            dy = self.f([y, t], training=training) # gradient
            y = y + (nt - t) * dy # next state
            res.append(y) # store result
        res = tf.stack(res, axis=1) # concatenate
        return res
```

- The `__call__` method implements the method using tensor operators



# Euler Method Model

The model for the Euler method is in the `ODEEulerModel` class

```
class ODEEulerModel(keras.Model):  
    ...  
  
    def train_step(self, data):  
        (y0, T), yt = data # unpack  
        with tf.GradientTape() as tape:  
            y = self.call([y0, T], training=True) # ODE integration  
            # Loss computation  
            mask = ~keras.ops.isnan(yt)  
            loss = self.compute_loss(x=None, y=yt[mask], y_pred=y[mask])  
        ...
```

- The loss is computed as usual on all available measurements
- We can exclude points by setting the corresponding target to NaN



# Training Set

## We have a single sequence of measurements

...Therefore, just a training set (no validation, no test)

- Our first input is the initial state:

```
In [8]: tr_y0 = np.array(data.iloc[0]).reshape(1, -1); display(tr_y0)
array([[0.]])
```

- The second is the sequence of evaluation points (time steps)

```
In [9]: tr_T = np.array(data.index).reshape(1, -1); display(tr_T[:, :30])
array([[ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10., 11., 12.,
        13., 14., 15., 16., 17., 18., 19., 20., 21., 22., 23., 24., 25.,
        26., 27., 28., 29.]])
```



# Training Set

## We have a single sequence of measurements

...Therefore, just a training set (no validation, no test)

- Then we need to prepare our ground truth

```
In [10]: tr_y = np.array(data['V']).reshape(1, -1)
tr_y[:, 0] = np.nan
display(tr_y[:, :30])
```

```
array([[ nan,  1.41003718,  2.6543906 ,  3.75252866,  4.7216321 ,
        5.57686288,  6.33160139,  6.99765581,  7.58544676,  8.10417046,
        8.56194251,  8.96592493,  9.32243816,  9.63705999,  9.91471279,
       10.15974045, 10.37597661, 10.56680439, 10.73520931, 10.88382613,
       11.01498002, 11.13072291, 11.23286566, 11.32300631, 11.40255515,
       11.47275676, 11.53470947, 11.58938254, 11.63763136, 11.6802108 ]])
```

- This is the sequence of all measurements, with the first state "masked"






# Training Process

We can now build and train the model

```
In [14]: %%time
dRC = util.RCNablaLayer(tau_ref=10, vs_ref=10)
euler = util.ODEEulerModel(dRC)
history = util.train_ml_model(euler, [tr_y0, tr_T], tr_y, validation_split=0.0, epochs=1000)
util.plot_training_history(history, figsize=figsize)
```



Model loss: 11.0890 (training)  
CPU times: user 7.96 s, sys: 2.36 s, total: 10.3 s

## Some Considerations

**It seems to be working! But there are a few issues**

First, the convergence is slow

- Stopping before ~500 epochs leads to less stable results

Second, we cannot use a validation set:

- This is due to the fact that we have a single sequence

Third, we are still **not getting** the correct parameters:

```
In [15]: print(f'tau: {tau:.2f} (real), {dRC.get_tau().numpy()[0]:.2f} (estimated)')
          print(f'Vs: {Vs:.2f} (real), {dRC.get_vs().numpy()[0]:.2f} (estimated)')
```

```
tau: 8.00 (real), 8.51 (estimated)
Vs: 12.00 (real), 12.00 (estimated)
```

**In the next section, we will see how to address these issues**

